



In general, you can see that the run times for the FA node selection are longer than the run times for MRV node selection. This is because the FA method just selects the first node with multiple domains that it encounters, not attempting to optimize the selection at all. This results in more subsequent iterations on average on account of the lack of selection optimization.

The MRV selection on the other hand picks nodes that are the most constrained, and thus produces the smallest number of subsequent iterations, which leads to faster runtimes in general. There is at least one case shown where the MRV runtime is longer than the FA runtime, denoted by a plot point below the diagonal line. This is a very particular case where the FA selection just happened to choose a faster path than MRV, and this does not happen in general.